

Implementation Plan -

AutoValuePredict ML

This document outlines the complete step-by-step implementation plan for the AutoValuePredict ML project, a machine learning system for predicting used car prices in Brazil.

Project Overview

Goal: Build an end-to-end ML pipeline that predicts the market value of used cars in Brazil, including data collection, preprocessing, feature engineering, model training, evaluation, and API deployment.

Current Status: Data collection and enrichment completed | EDA completed | Data preprocessing completed | Feature engineering completed | Baseline models completed | Advanced models completed | Model optimization completed | Model persistence completed | API development completed | Docker & Deployment next

Development Approach: MVP-first strategy - build a functional end-to-end pipeline with essential features, then iterate and optimize.

[!IMPORTANT] > **Data Limitations:** This project uses enriched FIPE data where features like `km` (mileage), `location`, `color`, `doors`, and `condition` are synthetically generated using statistical patterns. While realistic, these are not real-world observations. Model predictions should be validated against actual market data before production use.

[!NOTE] > **Timeline:** Estimated 12-16 weeks for complete implementation. Tasks are marked as **[ESSENTIAL]** (required for MVP) or **[OPTIONAL]** (enhancements). Focus on essential tasks first to achieve a working pipeline quickly.

Phase 1: Exploratory Data Analysis (EDA)

Completed

1.1 Initial Data Exploration

- [x] Load enriched datasets (`fipe_cars_enriched.csv` and `fipe_2022_enriched.csv`)
- [x] Basic data overview:
- [x] Dataset shape and memory usage
- [x] Column types and basic statistics
- [x] Missing values analysis
- [x] Duplicate records check
- [x] Create notebook: `notebooks/01_data_overview.ipynb`

1.2 Target Variable Analysis

- [x] Analyze price distribution:
- [x] Histogram and box plots
- [x] Skewness and kurtosis
- [x] Outlier detection (IQR method, Z-score)
- [x] Price ranges by vehicle category
- [x] Price trends over time (`year_of_reference`)
- [x] Price by brand, model, state
- [x] Create notebook: `notebooks/02_target_analysis.ipynb`

1.3 Feature Analysis

- [x] Categorical features:
- [x] Brand distribution and frequency
- [x] Model distribution
- [x] State/city distribution
- [x] Fuel type, transmission, color distributions
- [x] Condition distribution
- [x] Numerical features:
- [x] Year distribution and trends
- [x] Mileage (km) distribution and relationship with age
- [x] Engine size distribution

- [x] Doors distribution
- [x] Create notebook: `notebooks/03_feature_analysis.ipynb`

1.4 Relationships and Correlations

- [x] Correlation matrix (numerical features)
- [x] Price vs. age relationship
- [x] Price vs. mileage relationship
- [x] Price vs. brand/model analysis
- [x] Price vs. location (state) analysis
- [x] Feature interactions:
 - [x] Price by brand and year
 - [x] Price by fuel type and transmission
 - [x] Price by condition and age
- [x] Create notebook: `notebooks/04_correlations.ipynb`

1.5 Data Quality Assessment

- [x] Identify data quality issues:
- [x] Inconsistent values
- [x] Outliers that need treatment
- [x] Missing values (if any)
- [x] Data type issues
- [x] Document findings and recommendations
- [x] Create notebook: `notebooks/05_data_quality.ipynb`

Deliverables:

- 5 Jupyter notebooks with complete EDA
- Summary document with key findings
- Data quality report

Estimated Time: 1-2 weeks

Phase 2: Data Preprocessing & Cleaning

Completed

2.1 Data Cleaning Module

- [x] Create `src/data/cleaner.py`:
- [x] Remove duplicates
- [x] Handle outliers (price, km):
 - [x] IQR method for outliers
 - [x] Z-score method
 - [x] Domain knowledge-based limits
- [x] Handle missing values (if any)
- [x] Data type corrections
- [x] Standardize text fields (brand, model, city names)
- [x] Unit tests: `tests/test_cleaner.py`

2.2 Data Validation

- [x] Create `src/data/validator.py`:
- [x] Schema validation
- [x] Range checks (year, price, km)
- [x] Categorical value validation
- [x] Business rule validation
- [x] Unit tests: `tests/test_validator.py`

2.3 Data Splitting

- [x] Create `src/data/splitter.py`:
- [x] Train/validation/test split (70/15/15 or 80/10/10)
- [x] Stratified split by price ranges (optional)
- [x] Time-based split (if using `year_of_reference`)
- [x] Save splits to `data/processed/`
- [x] Unit tests: `tests/test_splitter.py`

Deliverables:

-  Data cleaning pipeline (`src/data/cleaner.py`)

- Data validation module (`src/data/validator.py`)
- Data splitting module (`src/data/splitter.py`)
- Modular ML pipeline system (`src/pipeline/`)
- Unit tests for all modules (`tests/test_*.py`)
- Validated and cleaned datasets
- Train/validation/test splits (70/15/15)
- Pipeline execution scripts (`scripts/run_pipeline.py`, `scripts/preprocess_data.py`)
- Makefile commands for pipeline execution

Estimated Time: 1 week

Actual Time: Completed

Results:

- Processed 747,948 rows (from 889,282 original)
 - Train: 523,563 rows (70%)
 - Validation: 112,192 rows (15%)
 - Test: 112,193 rows (15%)
 - All data quality checks passed
-

Phase 3: Feature Engineering Completed

3.1 Feature Creation - Phase 1 (Essential for MVP)

- [x] Create `src/features/engineering.py`:
- [x] **[ESSENTIAL] Basic temporal features:**
 - [x] Vehicle age (already exists, verify calculation)
 - [x] Age squared (non-linear relationship)
- [x] **[ESSENTIAL] Categorical encoding:**
 - [x] One-hot encoding for low cardinality features (fuel_type, transmission, condition)
 - [x] Target encoding for high cardinality (brand, model, state)
- [x] **[ESSENTIAL] Numerical transformations:**
 - [x] Log transformation for price (if skewed)
 - [x] Log transformation for km (if skewed)

- [x] Standardization/normalization
- [x] **[ESSENTIAL] Location features:**
 - [x] State encoding (target encoding)
 - [x] Region encoding (Norte, Nordeste, Sul, Sudeste, Centro-Oeste)

3.1.1 Feature Creation - Phase 2 (Optional Enhancements)

Implemented

- [x] **[OPTIONAL] Advanced features:**
- [x] Depreciation rate calculation
- [x] Frequency encoding (brand/model frequency)
- [x] **Interaction features:**
 - [x] Brand × Year
 - [x] Fuel × Transmission
 - [x] Age × Condition
 - [x] Km per year (km / age)
- [x] **Binning:**
 - [x] Price bins (for stratified splits)
 - [x] Age bins
 - [x] Mileage bins
- [x] **Advanced location features:**
 - [x] Region encoding (Norte, Nordeste, Sul, Sudeste, Centro-Oeste) -
Already in Phase 1
 - [x] City size category (if applicable)

Note: Advanced features are implemented in `AdvancedFeatureCreator` class and can be enabled via `use_advanced_features=True` parameter. They are disabled by default to maintain MVP focus.

3.2 Feature Selection

- [x] Create `src/features/selectors.py`:
- [x] Correlation-based feature selection
- [x] Mutual information
- [x] Feature importance from baseline models
- [x] Remove highly correlated features
- [x] Document selected features

3.3 Feature Pipeline

- [x] Create `src/features/pipeline.py`:
- [x] Scikit-learn Pipeline or custom pipeline
- [x] Combine all transformations
- [x] Fit on training, transform on validation/test
- [x] Save fitted pipeline for inference
- [x] Integrate FeatureEngineeringStep into main pipeline

Deliverables:

- Feature engineering pipeline (`src/features/pipeline.py`)
- Feature engineering modules (`src/features/engineering.py`, `src/features/selectors.py`)
- FeatureEngineeringStep integrated into main pipeline
- Pipeline persistence (save/load functionality)
- Engineered feature dataset (will be created when pipeline runs)
- Feature importance analysis (available when feature selection is enabled)

Estimated Time: 1-2 weeks

Actual Time: Completed

Phase 4: Baseline Models Completed

4.1 Baseline Implementations

- [x] Create `src/models/baseline.py`:
- [x] **Mean/Median baseline:** Simple average/median price
- [x] **Linear Regression:** Basic linear model
- [x] **Ridge Regression:** L2 regularization
- [x] **Lasso Regression:** L1 regularization
- [x] **Decision Tree:** Simple tree model
- [x] Evaluate all baselines
- [x] Document baseline performance

4.2 Evaluation Metrics

- [x] Create `src/models/evaluator.py`:
- [x] **RMSE** (Root Mean Squared Error)
- [x] **MAE** (Mean Absolute Error)
- [x] **MAPE** (Mean Absolute Percentage Error)
- [x] **R² Score** (Coefficient of Determination)
- [x] **Residual analysis:**
 - [x] Residual plots
 - [x] Q-Q plots
 - [x] Residual distribution
- [x] Create visualization functions for metrics

Deliverables:

- Baseline model implementations (`src/models/baseline.py`)
- Baseline performance report (saved to `models/baseline_results/`)
- Evaluation metrics module (`src/models/evaluator.py`)
- TrainBaselineModelsStep integrated into main pipeline
- Training script (`scripts/train_baseline_models.py`)
- Results saved to `models/baseline_results/` (separate from `data/processed/`)

Estimated Time: 3-5 days

Actual Time: Completed

Phase 5: Advanced Model Development

Completed

5.1 Model Implementations - Essential

- [x] Create `src/models/trainer.py`:
- [x] **[ESSENTIAL] Random Forest:**
 - [x] Hyperparameter tuning (RandomizedSearchCV with 5-fold CV)
 - [x] `n_estimators`, `max_depth`, `min_samples_split`, `min_samples_leaf`, `max_features`

- [x] **[ESSENTIAL] Gradient Boosting (XGBoost):**
 - [x] Hyperparameter tuning (validation-based search)
 - [x] learning_rate, n_estimators, max_depth, subsample, colsample_bytree
 - [x] Validation set monitoring (early stopping removed for compatibility)

5.1.1 Model Implementations - Optional

- [x] **[OPTIONAL] LightGBM:**
 - [x] Hyperparameter tuning (validation-based search)
 - [x] num_leaves, learning_rate, max_depth, feature_fraction
 - [x] Validation set monitoring (early stopping removed for compatibility)
- [] **[OPTIONAL] CatBoost:**
 - [] Good for categorical features
 - [] Hyperparameter tuning
 - **Note:** CatBoost not implemented yet (can be added in Phase 6 if needed)

5.2 Model Training

- [x] Implement cross-validation:
- [x] K-fold cross-validation (k=2) for Random Forest (reduced for memory efficiency)
- [x] Validation-based search for XGBoost and LightGBM
- [x] Train all models on training set
- [x] Validate on validation set
- [x] Track training time and model size

5.3 Model Comparison

- [x] Compare all models:
- [x] Performance metrics (RMSE, MAE, MAPE, R²)
- [x] Training time tracking
- [x] Validation scores tracking
- [x] Feature importance analysis:
- [x] Tree-based model feature importance (available via model attributes)
- [x] Create comparison report (CSV and visualization)

5.4 Model Selection

- [x] Select best model based on:
- [x] Validation performance (metrics comparison)
- [x] Training time tracking
- [] Final evaluation on test set (to be done in Phase 6)
- [] Document model selection rationale (to be done in Phase 6)

Deliverables:

- Advanced model trainer (`src/models/trainer.py`)
- TrainAdvancedModelsStep integrated into main pipeline
- Training script (`scripts/train_advanced_models.py`)
- Trained advanced models (saved to `models/advanced_results/`)
- Model comparison report (CSV and plots)
- Model artifacts saved (joblib format)
- Feature importance analysis (available via model attributes, detailed analysis in Phase 6)

Estimated Time: 2-3 weeks

Actual Time: Completed

Phase 6: Model Optimization & Fine-tuning In Progress

6.1 Hyperparameter Optimization

- [x] Fine-tune selected model:
- [x] GridSearchCV or RandomSearchCV (implemented in `scripts/optimize_lightgbm.py`)
- [] Bayesian optimization (Optuna) - optional (can be added later)
- [] Learning curves analysis (optional enhancement)
- [x] Optimize for business metric (MAPE or RMSE) - RMSE scorer implemented

6.2 Model Ensemble (Optional Enhancement)

- [] **[OPTIONAL]** Create `src/models/ensemble.py`:
- [] Voting regressor
- [] Stacking regressor
- [] Weighted average ensemble
- [] **[OPTIONAL]** Evaluate ensemble performance

[!TIP] Ensemble methods can improve performance by 2-5% but add complexity.
Consider only after achieving good results with single models.

6.3 Model Validation

- [x] Final test set evaluation (implemented in `scripts/evaluate_test_set.py`)
- [x] Performance by segments (implemented in `scripts/analyze_segments_and_errors.py`):
 - [x] By price range
 - [x] By brand
 - [x] By age
 - [x] By location (region)
- [x] Error analysis (implemented in `scripts/analyze_segments_and_errors.py`):
 - [x] Identify worst predictions
 - [x] Analyze error patterns
- [x] Document model limitations (via reports and visualizations)

Deliverables:

- Hyperparameter optimization script (`scripts/optimize_lightgbm.py`)
- Test set evaluation script (`scripts/evaluate_test_set.py`)
- Segment and error analysis script (`scripts/analyze_segments_and_errors.py`)
- Makefile commands for Phase 6 scripts
- Optimized model (to be generated when script runs)
- Final performance metrics (to be generated when scripts run)
- Model validation report (to be generated when scripts run)
- Error analysis (to be generated when scripts run)

Estimated Time: 1 week

Status: Scripts implemented  | Ready for execution 

Phase 7: Model Persistence & Versioning

Completed

7.1 Model Saving

- [x] Create `src/models/persistence.py`:
- [x] Save trained model (`joblib`)
- [x] Save feature pipeline
- [x] Save model metadata:
 - [x] Training date
 - [x] Performance metrics
 - [x] Feature list
 - [x] Hyperparameters
 - [x] Training info
- [x] Save to `models/` directory with versioned structure

7.2 Model Versioning

- [x] Implement model versioning:
- [x] Version naming convention (v1.0.0, v1.0.1, etc.)
- [x] Model registry (JSON file)
- [x] Model comparison tracking
- [x] Production model marking

7.3 Model Loading

- [x] Create model loading function
- [x] Validate loaded model
- [x] Test inference with loaded model
- [x] Scripts for saving and loading models

Deliverables:

- Model persistence module (`src/models/persistence.py`)
- Model versioning system with registry
- Model saving script (`scripts/save_model_with_versioning.py`)
- Model loading and validation script (`scripts/load_and_validate_model.py`)
- Makefile commands for model operations
- ModelMetadata and ModelPersistence classes

Estimated Time: 2-3 days

Actual Time: Completed

Phase 8: API Development Completed

8.1 FastAPI Application Structure

- [x] Create `src/api/main.py`:
- [x] FastAPI app initialization
- [x] Model loading on startup
- [x] Health check endpoint: `GET /health`
- [x] Model info endpoint: `GET /model/info`

8.2 Prediction Endpoints

- [x] Create `src/api/schemas.py`:
- [x] Pydantic models for request/response
- [x] Input validation schemas
- [x] Response schemas
- [x] Create `src/api/predictor.py`:
- [x] Single prediction endpoint: `POST /predict`
- [x] Batch prediction endpoint: `POST /predict/batch`
- [x] Input validation
- [x] Feature transformation
- [x] Model inference
- [x] Response formatting

8.3 Error Handling

- [x] Create `src/api/errors.py`:
- [x] Custom exception classes
- [x] Error handlers
- [x] Validation error responses
- [x] Model inference errors

8.4 API Documentation

- [x] Configure OpenAPI/Swagger documentation
- [x] Add endpoint descriptions
- [x] Add example requests/responses
- [x] Document error codes

8.5 API Testing

- [x] Create `tests/test_api.py`:
- [x] Test health endpoint
- [x] Test prediction endpoints
- [x] Test input validation
- [x] Test error handling
- [x] Integration tests

Deliverables:

- FastAPI application (`src/api/main.py`)
- Prediction endpoints (`POST /predict`, `POST /predict/batch`)
- API documentation (OpenAPI/Swagger at `/docs`)
- API tests (`tests/test_api.py`)
- Error handling (`src/api/errors.py`)
- Request/Response schemas (`src/api/schemas.py`)
- Prediction logic (`src/api/predictor.py`)
- Makefile commands for API (`api-start`, `api-logs`, `api-health`, `test-api`)

Estimated Time: 1-2 weeks

Actual Time: Completed

Phase 9: Docker & Deployment

9.1 Docker Configuration Review

- [] Review and optimize Dockerfile:
- [] Multi-stage build (if needed)
- [] Minimize image size
- [] Optimize layer caching
- [] Review docker-compose.yml
- [] Test Docker build locally

9.2 Environment Configuration

- [] Create `.env.example`:
- [] API configuration
- [] Model path
- [] Logging level
- [] Update docker-compose.yml for environment variables

9.3 Deployment Preparation

- [] Create deployment documentation
- [] Test containerized application
- [] Performance testing:
- [] Load testing (optional)
- [] Response time testing
- [] Resource requirements documentation

9.4 CI/CD (Optional Enhancement)

- [] **[OPTIONAL]** Set up GitHub Actions:
- [] Run tests on push
- [] Code quality checks (black, flake8)
- [] Build Docker image
- [] Deploy to staging (optional)

[!NOTE] CI/CD is valuable for production systems but not essential for MVP.
Consider implementing after core functionality is complete.

Deliverables:

- Optimized Docker configuration
- Deployment documentation
- CI/CD pipeline (optional)

Estimated Time: 1 week

Phase 10: Documentation & Testing

10.1 Code Documentation

- [] Add docstrings to all functions and classes
- [] Document module purposes
- [] Add type hints throughout codebase
- [] Create API documentation

10.2 Project Documentation

- [] Update README.md with:
- [] Complete setup instructions
- [] Usage examples
- [] API usage guide
- [] Model information
- [] Create CONTRIBUTING.md (if applicable)
- [] Create ARCHITECTURE.md:
- [] Project structure
- [] Data flow
- [] Model architecture

10.3 Testing

- [] Unit tests for all modules:
- [] Data processing tests
- [] Feature engineering tests
- [] Model training tests
- [] API tests

- [] Integration tests
- [] Achieve >80% test coverage
- [] Create `tests/README.md` with testing instructions

10.4 Usage Examples

- [] Create example notebooks:
- [] Model training example
- [] API usage example
- [] Prediction example
- [] Create example scripts

Deliverables:

- Complete code documentation
- Updated project documentation
- Comprehensive test suite
- Usage examples

Estimated Time: 1-2 weeks

Phase 11: Optimization & Refinement

11.1 Performance Optimization

- [] Profile code for bottlenecks
- [] Optimize data loading
- [] Optimize feature engineering pipeline
- [] Optimize model inference:
 - [] Batch processing
 - [] Model quantization (optional)
 - [] Caching strategies

11.2 Code Quality

- [] Code review and refactoring
- [] Follow PEP 8 style guide

- [] Remove unused code
- [] Improve error messages
- [] Add logging throughout application

11.3 Model Monitoring (Optional Enhancement)

- [] [OPTIONAL] Create monitoring dashboard
- [] [OPTIONAL] Track prediction distribution
- [] [OPTIONAL] Monitor model drift
- [] [OPTIONAL] Set up alerts

[!WARNING] Model monitoring is critical for production systems but requires additional infrastructure. For portfolio/demo purposes, focus on core ML pipeline first.

Deliverables:

- Optimized codebase
- Performance improvements
- Code quality improvements

Estimated Time: 1 week

Implementation Timeline

[!NOTE] > **MVP Strategy:** Focus on essential tasks first to build a working end-to-end pipeline (Phases 1-8 core features). Optional enhancements can be added iteratively.

Phase	Task	Estimated Time	Priority	Status
1	Exploratory Data Analysis	1-2 weeks	Essential	<input checked="" type="checkbox"/> Completed
2	Data Preprocessing & Cleaning	1 week	Essential	<input checked="" type="checkbox"/> Completed
3		1 week	Essential	<input checked="" type="checkbox"/> Completed

Phase	Task	Estimated Time	Priority	Status
	Feature Engineering (Essential)			
3.1	Feature Engineering (Optional)	1 week	Optional	Completed
4	Baseline Models	3-5 days	Essential	Completed
5	Advanced Models (RF + XGBoost)	1-2 weeks	Essential	Completed
5.1	Additional Models (LightGBM, CatBoost)	1 week	Optional	Completed (LightGBM)
6	Model Optimization	1 week	Essential	Completed
6.1	Model Ensemble	3-5 days	Optional	Pending
7	Model Persistence	2-3 days	Essential	Completed
8	API Development	1-2 weeks	Essential	Completed
9	Docker & Deployment	1 week	Essential	Pending
9.1	CI/CD Pipeline	3-5 days	Optional	Pending
10	Documentation & Testing	1-2 weeks	Essential	Pending
11	Optimization & Refinement	1 week	Essential	Pending
11.1	Model Monitoring	3-5 days	Optional	Pending

MVP Timeline (Essential Only): 10-12 weeks

Full Implementation (Essential + Optional): 12-16 weeks

Key Deliverables Checklist

Data & Analysis

- [x] Raw datasets collected
- [x] Data enrichment completed
- [x] EDA notebooks completed
- [x] Data cleaning pipeline
- [x] Feature engineering pipeline

Models

- [x] Baseline models implemented
- [x] Advanced models trained (Random Forest, XGBoost, LightGBM)
- [] Best model selected and optimized (to be done in Phase 6)
- [x] Model artifacts saved (baseline results in `models/baseline_results/`, advanced results in `models/advanced_results/`)
- [x] Model performance report (baseline and advanced models)

API & Deployment

- [x] FastAPI application
- [x] Prediction endpoints
- [] Docker configuration
- [] Deployment documentation

Documentation & Quality

- [] Complete code documentation
 - [] Project documentation
 - [] Test suite (>80% coverage)
 - [] Usage examples
-

Success Criteria

- [] Model achieves acceptable performance:

- [] MAPE < 15-20% (or RMSE < acceptable threshold)
 - [] R² > 0.85
 - [] API responds to predictions in < 1 second
 - [] Code is well-documented and tested (>80% coverage)
 - [] Project can be easily reproduced and deployed
 - [] All phases are completed and documented
-

Notes & Considerations

Development Best Practices

- Data Quality:** Continuously monitor data quality throughout the pipeline
- Reproducibility:** Use random seeds for all random operations
- Version Control:** Track model versions and data versions
- Performance:** Balance model accuracy with inference speed
- Interpretability:** Consider model interpretability for business stakeholders
- Scalability:** Design pipeline to handle larger datasets in the future

Production Readiness (Future Considerations)

[!WARNING] > Synthetic Data Validation: Before deploying to production, validate model predictions against real-world market data. The current dataset uses synthetic features that may not capture all market dynamics.

- Model Retraining:** Plan for periodic model retraining (monthly/quarterly) as car market conditions change
 - Drift Monitoring:** Implement data drift detection to identify when model performance degrades
 - A/B Testing:** Consider A/B testing framework for comparing model versions in production
 - Business Validation:** Validate predictions with automotive market experts before production deployment
 - Edge Cases:** Document and handle edge cases (luxury cars, rare models, extreme mileage)
-

Next Steps

1. **Start with Phase 1:** Begin Exploratory Data Analysis
 2. **Create first notebook:** `notebooks/01_data_overview.ipynb`
 3. **Set up development environment:** Ensure all dependencies are installed
 4. **Review data:** Load and inspect enriched datasets
-

Last Updated: 2025-01-26

Current Phase: Phase 8 - API Development  Completed | Phase 9 - Docker & Deployment  Next

Strategy: MVP-first approach with essential features, then iterate with optional enhancements