

# Implementation Plan - AutoValuePredict ML

---

This document outlines the complete step-by-step implementation plan for the AutoValuePredict ML project, a machine learning system for predicting used car prices in Brazil.

## Project Overview

---

**Goal:** Build an end-to-end ML pipeline that predicts the market value of used cars in Brazil, including data collection, preprocessing, feature engineering, model training, evaluation, and API deployment.

**Current Status:** ✓ Data collection and enrichment completed | ✓ EDA completed | ✓ Data preprocessing completed | ✓ Feature engineering completed | ✓ Baseline models completed | 🚧 Advanced models next

**Development Approach:** MVP-first strategy - build a functional end-to-end pipeline with essential features, then iterate and optimize.

[!IMPORTANT] > **Data Limitations:** This project uses enriched FIPE data where features like `km` (mileage), `location`, `color`, `doors`, and `condition` are synthetically generated using statistical patterns. While realistic, these are not real-world observations. Model predictions should be validated against actual market data before production use.

[!NOTE] > **Timeline:** Estimated 12-16 weeks for complete implementation. Tasks are marked as **[ESSENTIAL]** (required for MVP) or **[OPTIONAL]** (enhancements). Focus on essential tasks first to achieve a working pipeline quickly.

---

# Phase 1: Exploratory Data Analysis (EDA)

## Completed

---

### 1.1 Initial Data Exploration

- [x] Load enriched datasets (`fipe_cars_enriched.csv` and `fipe_2022_enriched.csv`)
- [x] Basic data overview:
- [x] Dataset shape and memory usage
- [x] Column types and basic statistics
- [x] Missing values analysis
- [x] Duplicate records check
- [x] Create notebook: `notebooks/01_data_overview.ipynb`

### 1.2 Target Variable Analysis

- [x] Analyze price distribution:
- [x] Histogram and box plots
- [x] Skewness and kurtosis
- [x] Outlier detection (IQR method, Z-score)
- [x] Price ranges by vehicle category
- [x] Price trends over time (`year_of_reference`)
- [x] Price by brand, model, state
- [x] Create notebook: `notebooks/02_target_analysis.ipynb`

### 1.3 Feature Analysis

- [x] Categorical features:
- [x] Brand distribution and frequency
- [x] Model distribution
- [x] State/city distribution
- [x] Fuel type, transmission, color distributions
- [x] Condition distribution
- [x] Numerical features:
- [x] Year distribution and trends
- [x] Mileage (km) distribution and relationship with age
- [x] Engine size distribution

- [x] Doors distribution
- [x] Create notebook: `notebooks/03_feature_analysis.ipynb`

## 1.4 Relationships and Correlations

- [x] Correlation matrix (numerical features)
- [x] Price vs. age relationship
- [x] Price vs. mileage relationship
- [x] Price vs. brand/model analysis
- [x] Price vs. location (state) analysis
- [x] Feature interactions:
  - [x] Price by brand and year
  - [x] Price by fuel type and transmission
  - [x] Price by condition and age
- [x] Create notebook: `notebooks/04_correlations.ipynb`

## 1.5 Data Quality Assessment

- [x] Identify data quality issues:
- [x] Inconsistent values
- [x] Outliers that need treatment
- [x] Missing values (if any)
- [x] Data type issues
- [x] Document findings and recommendations
- [x] Create notebook: `notebooks/05_data_quality.ipynb`

### **Deliverables:**

- 5 Jupyter notebooks with complete EDA
- Summary document with key findings
- Data quality report

**Estimated Time:** 1-2 weeks

---

# Phase 2: Data Preprocessing & Cleaning

## Completed

---

### 2.1 Data Cleaning Module

- [x] Create `src/data/cleaner.py`:
- [x] Remove duplicates
- [x] Handle outliers (price, km):
  - [x] IQR method for outliers
  - [x] Z-score method
  - [x] Domain knowledge-based limits
- [x] Handle missing values (if any)
- [x] Data type corrections
- [x] Standardize text fields (brand, model, city names)
- [x] Unit tests: `tests/test_cleaner.py`

### 2.2 Data Validation

- [x] Create `src/data/validator.py`:
- [x] Schema validation
- [x] Range checks (year, price, km)
- [x] Categorical value validation
- [x] Business rule validation
- [x] Unit tests: `tests/test_validator.py`

### 2.3 Data Splitting

- [x] Create `src/data/splitter.py`:
- [x] Train/validation/test split (70/15/15 or 80/10/10)
- [x] Stratified split by price ranges (optional)
- [x] Time-based split (if using `year_of_reference`)
- [x] Save splits to `data/processed/`
- [x] Unit tests: `tests/test_splitter.py`

#### Deliverables:

-  Data cleaning pipeline (`src/data/cleaner.py`)

- Data validation module (`src/data/validator.py`)
- Data splitting module (`src/data/splitter.py`)
- Modular ML pipeline system (`src/pipeline/`)
- Unit tests for all modules (`tests/test_*.py`)
- Validated and cleaned datasets
- Train/validation/test splits (70/15/15)
- Pipeline execution scripts (`scripts/run_pipeline.py`, `scripts/preprocess_data.py`)
- Makefile commands for pipeline execution

**Estimated Time:** 1 week

**Actual Time:** Completed

#### Results:

- Processed 747,948 rows (from 889,282 original)
  - Train: 523,563 rows (70%)
  - Validation: 112,192 rows (15%)
  - Test: 112,193 rows (15%)
  - All data quality checks passed
- 

## Phase 3: Feature Engineering Completed

### 3.1 Feature Creation - Phase 1 (Essential for MVP)

- [x] Create `src/features/engineering.py`:
- [x] **[ESSENTIAL] Basic temporal features:**
  - [x] Vehicle age (already exists, verify calculation)
  - [x] Age squared (non-linear relationship)
- [x] **[ESSENTIAL] Categorical encoding:**
  - [x] One-hot encoding for low cardinality features (fuel\_type, transmission, condition)
  - [x] Target encoding for high cardinality (brand, model, state)
- [x] **[ESSENTIAL] Numerical transformations:**
  - [x] Log transformation for price (if skewed)
  - [x] Log transformation for km (if skewed)

- [x] Standardization/normalization
- [x] **[ESSENTIAL] Location features:**
  - [x] State encoding (target encoding)
  - [x] Region encoding (Norte, Nordeste, Sul, Sudeste, Centro-Oeste)

### 3.1.1 Feature Creation - Phase 2 (Optional Enhancements)

#### Implemented

- [x] **[OPTIONAL] Advanced features:**
  - [x] Depreciation rate calculation
  - [x] Frequency encoding (brand/model frequency)
- [x] **Interaction features:**
  - [x] Brand × Year
  - [x] Fuel × Transmission
  - [x] Age × Condition
  - [x] Km per year (km / age)
- [x] **Binning:**
  - [x] Price bins (for stratified splits)
  - [x] Age bins
  - [x] Mileage bins
- [x] **Advanced location features:**
  - [x] Region encoding (Norte, Nordeste, Sul, Sudeste, Centro-Oeste) -  
Already in Phase 1
  - [x] City size category (if applicable)

**Note:** Advanced features are implemented in `AdvancedFeatureCreator` class and can be enabled via `use_advanced_features=True` parameter. They are disabled by default to maintain MVP focus.

## 3.2 Feature Selection

- [x] Create `src/features/selectors.py`:
- [x] Correlation-based feature selection
- [x] Mutual information
- [x] Feature importance from baseline models
- [x] Remove highly correlated features
- [x] Document selected features

### 3.3 Feature Pipeline

- [x] Create `src/features/pipeline.py`:
- [x] Scikit-learn Pipeline or custom pipeline
- [x] Combine all transformations
- [x] Fit on training, transform on validation/test
- [x] Save fitted pipeline for inference
- [x] Integrate FeatureEngineeringStep into main pipeline

#### Deliverables:

- Feature engineering pipeline (`src/features/pipeline.py`)
- Feature engineering modules (`src/features/engineering.py`, `src/features/selectors.py`)
- FeatureEngineeringStep integrated into main pipeline
- Pipeline persistence (save/load functionality)
- Engineered feature dataset (will be created when pipeline runs)
- Feature importance analysis (available when feature selection is enabled)

**Estimated Time:** 1-2 weeks

**Actual Time:** Completed

---

## Phase 4: Baseline Models Completed

---

### 4.1 Baseline Implementations

- [x] Create `src/models/baseline.py`:
- [x] **Mean/Median baseline:** Simple average/median price
- [x] **Linear Regression:** Basic linear model
- [x] **Ridge Regression:** L2 regularization
- [x] **Lasso Regression:** L1 regularization
- [x] **Decision Tree:** Simple tree model
- [x] Evaluate all baselines
- [x] Document baseline performance

## 4.2 Evaluation Metrics

- [x] Create `src/models/evaluator.py`:
- [x] **RMSE** (Root Mean Squared Error)
- [x] **MAE** (Mean Absolute Error)
- [x] **MAPE** (Mean Absolute Percentage Error)
- [x] **R<sup>2</sup> Score** (Coefficient of Determination)
- [x] **Residual analysis:**
  - [x] Residual plots
  - [x] Q-Q plots
  - [x] Residual distribution
- [x] Create visualization functions for metrics

### Deliverables:

- Baseline model implementations (`src/models/baseline.py`)
- Baseline performance report (saved to `models/baseline_results/`)
- Evaluation metrics module (`src/models/evaluator.py`)
- TrainBaselineModelsStep integrated into main pipeline
- Training script (`scripts/train_baseline_models.py`)
- Results saved to `models/baseline_results/` (separate from `data/processed/`)

**Estimated Time:** 3-5 days

**Actual Time:** Completed

---

## Phase 5: Advanced Model Development

---

### 5.1 Model Implementations - Essential

- [ ] Create `src/models/trainer.py`:
- [ ] **[ESSENTIAL] Random Forest:**
  - [ ] Hyperparameter tuning (GridSearchCV/RandomSearchCV)
  - [ ] n\_estimators, max\_depth, min\_samples\_split
- [ ] **[ESSENTIAL] Gradient Boosting (XGBoost):**
  - [ ] Hyperparameter tuning
  - [ ] learning\_rate, n\_estimators, max\_depth

- [ ] Early stopping

### 5.1.1 Model Implementations - Optional

- [ ] **[OPTIONAL] LightGBM:**
- [ ] Hyperparameter tuning
- [ ] num\_leaves, learning\_rate, feature\_fraction
- [ ] Early stopping
- [ ] **[OPTIONAL] CatBoost:**
- [ ] Good for categorical features
- [ ] Hyperparameter tuning

## 5.2 Model Training

- [ ] Implement cross-validation:
- [ ] K-fold cross-validation (k=5)
- [ ] Time-based cross-validation (if applicable)
- [ ] Train all models on training set
- [ ] Validate on validation set
- [ ] Track training time and model size

## 5.3 Model Comparison

- [ ] Compare all models:
- [ ] Performance metrics (RMSE, MAE, MAPE, R<sup>2</sup>)
- [ ] Training time
- [ ] Inference time
- [ ] Model complexity
- [ ] Feature importance analysis:
- [ ] Tree-based model feature importance
- [ ] Permutation importance
- [ ] Create comparison report

## 5.4 Model Selection

- [ ] Select best model based on:
- [ ] Validation performance
- [ ] Generalization (cross-validation)

- [ ] Inference speed
- [ ] Model interpretability
- [ ] Final evaluation on test set
- [ ] Document model selection rationale

**Deliverables:**

- Trained advanced models
- Model comparison report
- Selected best model
- Feature importance analysis

**Estimated Time:** 2-3 weeks

---

## Phase 6: Model Optimization & Fine-tuning

---

### 6.1 Hyperparameter Optimization

- [ ] Fine-tune selected model:
- [ ] GridSearchCV or RandomSearchCV
- [ ] Bayesian optimization (Optuna) - optional
- [ ] Learning curves analysis
- [ ] Optimize for business metric (MAPE or RMSE)

### 6.2 Model Ensemble (Optional Enhancement)

- [ ] **[OPTIONAL]** Create `src/models/ensemble.py`:
- [ ] Voting regressor
- [ ] Stacking regressor
- [ ] Weighted average ensemble
- [ ] **[OPTIONAL]** Evaluate ensemble performance

| **[!TIP]** Ensemble methods can improve performance by 2-5% but add complexity.  
Consider only after achieving good results with single models.

## 6.3 Model Validation

- [ ] Final test set evaluation
- [ ] Performance by segments:
- [ ] By price range
- [ ] By brand
- [ ] By age
- [ ] By location
- [ ] Error analysis:
- [ ] Identify worst predictions
- [ ] Analyze error patterns
- [ ] Document model limitations

### **Deliverables:**

- Optimized model
- Final performance metrics
- Model validation report
- Error analysis

**Estimated Time:** 1 week

---

## Phase 7: Model Persistence & Versioning

---

### 7.1 Model Saving

- [ ] Create `src/models/persistence.py`:
- [ ] Save trained model (joblib/pickle)
- [ ] Save feature pipeline
- [ ] Save preprocessing steps
- [ ] Save model metadata:
  - [ ] Training date
  - [ ] Performance metrics
  - [ ] Feature list
  - [ ] Hyperparameters
- [ ] Save to `models/` directory

## 7.2 Model Versioning

- [ ] Implement model versioning:
- [ ] Version naming convention (e.g., v1.0.0)
- [ ] Model registry (simple JSON file or MLflow)
- [ ] Model comparison tracking

## 7.3 Model Loading

- [ ] Create model loading function
- [ ] Validate loaded model
- [ ] Test inference with loaded model

### **Deliverables:**

- Model persistence module
- Saved model artifacts
- Model versioning system

**Estimated Time:** 2-3 days

---

# Phase 8: API Development

---

## 8.1 FastAPI Application Structure

- [ ] Create `src/api/main.py`:
- [ ] FastAPI app initialization
- [ ] Model loading on startup
- [ ] Health check endpoint: `GET /health`
- [ ] Model info endpoint: `GET /model/info`

## 8.2 Prediction Endpoints

- [ ] Create `src/api/schemas.py`:
- [ ] Pydantic models for request/response
- [ ] Input validation schemas
- [ ] Response schemas

- [ ] Create `src/api/predictor.py`:
- [ ] Single prediction endpoint: `POST /predict`
- [ ] Batch prediction endpoint: `POST /predict/batch`
- [ ] Input validation
- [ ] Feature transformation
- [ ] Model inference
- [ ] Response formatting

### 8.3 Error Handling

- [ ] Create `src/api/errors.py`:
- [ ] Custom exception classes
- [ ] Error handlers
- [ ] Validation error responses
- [ ] Model inference errors

### 8.4 API Documentation

- [ ] Configure OpenAPI/Swagger documentation
- [ ] Add endpoint descriptions
- [ ] Add example requests/responses
- [ ] Document error codes

### 8.5 API Testing

- [ ] Create `tests/test_api.py`:
- [ ] Test health endpoint
- [ ] Test prediction endpoints
- [ ] Test input validation
- [ ] Test error handling
- [ ] Integration tests

#### **Deliverables:**

- FastAPI application
- Prediction endpoints
- API documentation
- API tests

**Estimated Time:** 1-2 weeks

---

## Phase 9: Docker & Deployment

---

### 9.1 Docker Configuration Review

- [ ] Review and optimize Dockerfile:
- [ ] Multi-stage build (if needed)
- [ ] Minimize image size
- [ ] Optimize layer caching
- [ ] Review docker-compose.yml
- [ ] Test Docker build locally

### 9.2 Environment Configuration

- [ ] Create `.env.example`:
- [ ] API configuration
- [ ] Model path
- [ ] Logging level
- [ ] Update docker-compose.yml for environment variables

### 9.3 Deployment Preparation

- [ ] Create deployment documentation
- [ ] Test containerized application
- [ ] Performance testing:
  - [ ] Load testing (optional)
  - [ ] Response time testing
- [ ] Resource requirements documentation

### 9.4 CI/CD (Optional Enhancement)

- [ ] **[OPTIONAL]** Set up GitHub Actions:
  - [ ] Run tests on push
  - [ ] Code quality checks (black, flake8)
  - [ ] Build Docker image

- [ ] Deploy to staging (optional)

[!NOTE] CI/CD is valuable for production systems but not essential for MVP.  
Consider implementing after core functionality is complete.

**Deliverables:**

- Optimized Docker configuration
- Deployment documentation
- CI/CD pipeline (optional)

**Estimated Time:** 1 week

---

## Phase 10: Documentation & Testing

---

### 10.1 Code Documentation

- [ ] Add docstrings to all functions and classes
- [ ] Document module purposes
- [ ] Add type hints throughout codebase
- [ ] Create API documentation

### 10.2 Project Documentation

- [ ] Update README.md with:
- [ ] Complete setup instructions
- [ ] Usage examples
- [ ] API usage guide
- [ ] Model information
- [ ] Create CONTRIBUTING.md (if applicable)
- [ ] Create ARCHITECTURE.md:
- [ ] Project structure
- [ ] Data flow
- [ ] Model architecture

## 10.3 Testing

- [ ] Unit tests for all modules:
- [ ] Data processing tests
- [ ] Feature engineering tests
- [ ] Model training tests
- [ ] API tests
- [ ] Integration tests
- [ ] Achieve >80% test coverage
- [ ] Create `tests/README.md` with testing instructions

## 10.4 Usage Examples

- [ ] Create example notebooks:
- [ ] Model training example
- [ ] API usage example
- [ ] Prediction example
- [ ] Create example scripts

### **Deliverables:**

- Complete code documentation
- Updated project documentation
- Comprehensive test suite
- Usage examples

**Estimated Time:** 1-2 weeks

---

# Phase 11: Optimization & Refinement

---

## 11.1 Performance Optimization

- [ ] Profile code for bottlenecks
- [ ] Optimize data loading
- [ ] Optimize feature engineering pipeline
- [ ] Optimize model inference:
- [ ] Batch processing

- [ ] Model quantization (optional)
- [ ] Caching strategies

## 11.2 Code Quality

- [ ] Code review and refactoring
- [ ] Follow PEP 8 style guide
- [ ] Remove unused code
- [ ] Improve error messages
- [ ] Add logging throughout application

## 11.3 Model Monitoring (Optional Enhancement)

- [ ] **[OPTIONAL]** Create monitoring dashboard
- [ ] **[OPTIONAL]** Track prediction distribution
- [ ] **[OPTIONAL]** Monitor model drift
- [ ] **[OPTIONAL]** Set up alerts

[!WARNING] Model monitoring is critical for production systems but requires additional infrastructure. For portfolio/demo purposes, focus on core ML pipeline first.

### **Deliverables:**

- Optimized codebase
- Performance improvements
- Code quality improvements

**Estimated Time:** 1 week

---

## Implementation Timeline

[!NOTE] > **MVP Strategy:** Focus on essential tasks first to build a working end-to-end pipeline (Phases 1-8 core features). Optional enhancements can be added iteratively.

Phase	Task	Estimated Time	Priority	Status
1	Exploratory Data Analysis	1-2 weeks	Essential	<span>✓</span> Completed
2	Data Preprocessing & Cleaning	1 week	Essential	<span>✓</span> Completed
3	Feature Engineering (Essential)	1 week	Essential	<span>✓</span> Completed
3.1	Feature Engineering (Optional)	1 week	Optional	<span>✓</span> Completed
4	Baseline Models	3-5 days	Essential	<span>✓</span> Completed
5	Advanced Models (RF + XGBoost)	1-2 weeks	Essential	<span>⌚</span> Pending
5.1	Additional Models (LightGBM, CatBoost)	1 week	Optional	<span>⌚</span> Pending
6	Model Optimization	1 week	Essential	<span>⌚</span> Pending
6.1	Model Ensemble	3-5 days	Optional	<span>⌚</span> Pending
7	Model Persistence	2-3 days	Essential	<span>⌚</span> Pending
8	API Development	1-2 weeks	Essential	<span>⌚</span> Pending
9	Docker & Deployment	1 week	Essential	<span>⌚</span> Pending
9.1	CI/CD Pipeline	3-5 days	Optional	<span>⌚</span> Pending
10	Documentation & Testing	1-2 weeks	Essential	<span>⌚</span> Pending
11	Optimization & Refinement	1 week	Essential	<span>⌚</span> Pending
11.1	Model Monitoring	3-5 days	Optional	<span>⌚</span> Pending

**MVP Timeline (Essential Only):** 10-12 weeks

**Full Implementation (Essential + Optional):** 12-16 weeks

---

## Key Deliverables Checklist

---

### Data & Analysis

- [x] Raw datasets collected
- [x] Data enrichment completed
- [x] EDA notebooks completed
- [x] Data cleaning pipeline
- [x] Feature engineering pipeline

### Models

- [x] Baseline models implemented
- [ ] Advanced models trained
- [ ] Best model selected and optimized
- [x] Model artifacts saved (baseline results in `models/baseline_results/`)
- [x] Model performance report (baseline models)

### API & Deployment

- [ ] FastAPI application
- [ ] Prediction endpoints
- [ ] Docker configuration
- [ ] Deployment documentation

### Documentation & Quality

- [ ] Complete code documentation
  - [ ] Project documentation
  - [ ] Test suite (>80% coverage)
  - [ ] Usage examples
-

# Success Criteria

---

- [ ] Model achieves acceptable performance:
  - [ ] MAPE < 15-20% (or RMSE < acceptable threshold)
  - [ ] R<sup>2</sup> > 0.85
  - [ ] API responds to predictions in < 1 second
  - [ ] Code is well-documented and tested (>80% coverage)
  - [ ] Project can be easily reproduced and deployed
  - [ ] All phases are completed and documented
- 

## Notes & Considerations

---

### Development Best Practices

1. **Data Quality:** Continuously monitor data quality throughout the pipeline
2. **Reproducibility:** Use random seeds for all random operations
3. **Version Control:** Track model versions and data versions
4. **Performance:** Balance model accuracy with inference speed
5. **Interpretability:** Consider model interpretability for business stakeholders
6. **Scalability:** Design pipeline to handle larger datasets in the future

### Production Readiness (Future Considerations)

[!WARNING] > **Synthetic Data Validation:** Before deploying to production, validate model predictions against real-world market data. The current dataset uses synthetic features that may not capture all market dynamics.

1. **Model Retraining:** Plan for periodic model retraining (monthly/quarterly) as car market conditions change
2. **Drift Monitoring:** Implement data drift detection to identify when model performance degrades
3. **A/B Testing:** Consider A/B testing framework for comparing model versions in production
4. **Business Validation:** Validate predictions with automotive market experts before production deployment

5. **Edge Cases:** Document and handle edge cases (luxury cars, rare models, extreme mileage)
- 

## Next Steps

---

1. **Start with Phase 1:** Begin Exploratory Data Analysis
  2. **Create first notebook:** notebooks/01\_data\_overview.ipynb
  3. **Set up development environment:** Ensure all dependencies are installed
  4. **Review data:** Load and inspect enriched datasets
- 

**Last Updated:** 2024-12-08

**Current Phase:** Phase 4 - Baseline Models  Completed | Phase 5 - Advanced

Models  Next

**Strategy:** MVP-first approach with essential features, then iterate with optional enhancements